

# Silicon Core Harvesting is Visible in Supercomputer Telemetry: A Fleet-Scale Study of 1,960 POWER9 Sockets

Mohsen Seyedkazemi Ardebili

*Affiliation*

mohsen.seyedkazemi@gmail.com

[TODO: co-authors]

**Abstract**—Processor vendors sell partially defective dies by permanently disabling faulty units, a practice known as core harvesting. The resulting per-die disable maps are proprietary and are conventionally invisible to system owners. We show that the out-of-band (BMC/IPMI) sensor telemetry routinely collected by production supercomputers exposes these maps directly, and we recover them for all 1,960 processor sockets of CINECA’s Marconi100 from the public M100 ExaData release [1]. Every socket exposes exactly 16 of 24 possible per-core thermal sensors, and the eight absent sensors always form four complete *slice* pairs  $(2k, 2k+1)$ —the POWER9 unit that shares an L2/L3 block [2], [3]—with zero exceptions. We observe 442 of the 495 combinatorially possible patterns, establishing that the map is per-die rather than a fixed SKU. Using a curveball null that preserves both per-socket and per-slice marginals, we show harvested slices cluster in sensor-index space ( $z = -20.8$ ). Thermal analysis confirms that the two cores of a slice are physically abutted (residual correlation  $r = +0.51$  within-slice versus  $+0.29$  for cross-slice neighbours at identical index distance), but we find only weak evidence that the slice index itself is monotonic in physical die position ( $p = 0.025$ ; a  $3 \times 4$  grid fits the thermal coupling better than a linear arrangement). Finally, we detect a sharp procurement-lot boundary at rack 22 (Welch  $t = 27.3$ ) separating two populations with markedly different harvesting statistics.

**Index Terms**—core harvesting, silicon binning, HPC telemetry, POWER9, yield, hardware provenance

## I. INTRODUCTION

[TODO: write]. Framing: (i) core harvesting is standard industry practice but per-die maps are undocumented; (ii) HPC operators now collect holistic out-of-band telemetry [1], which we show carries an unintended side channel revealing manufacturing outcomes; (iii) contributions are a recovery method, the first fleet-scale characterisation of harvesting in a deployed Top-10 system, evidence of index-space clustering and of procurement-lot structure, and a thermal method for testing die-topology hypotheses.

## II. BACKGROUND

### A. POWER9 slice organisation

The POWER9 die used in the IBM Power System AC922 contains 24 SMT4 cores (equivalently 12 SMT8 cores) on a  $695 \text{ mm}^2$  die in 14 nm FinFET with 17 metal layers [3]. Critically, *each pair of SMT4 cores—or one singleton SMT8 core—comprises a slice, and each slice owns 512 kB of L2*

*and 10 MB of L3* [3], [2]. IBM’s Hot Chips 28 presentation shows the 24 cores above  $12 \times 10 \text{ MB}$  L3 blocks, attached to a 7 TB/s on-chip fabric at  $256 \text{ GB/s} \times 12$  [2]. The slice is therefore the natural fusing granularity: disabling one SMT4 core would strand its half of a shared L2/L3 block.

Marconi100 nodes are AC922 model 8335-GTG with  $2 \times 16$ -core POWER9 at 2.6/3.1 GHz and  $4 \times$  NVIDIA V100 [4], [5]. **A 16-core part is a 24-core die with four slices fused off**; this is the object of our study. Fig. 1 shows the organisation annotated with the harvest rates we measure.

### B. The M100 ExaData dataset

M100 ExaData [1] publishes 49.9 TB of telemetry from Marconi100 covering 2020-03 to 2022-09 across 980+ nodes. The IPMI plugin collects BMC sensor data at 20 s cadence. The documentation describes the core-temperature family as `pX_coreY_temp`, “Temperature of core n. Y in the CPU socket n. X. X=0..1, Y=0..23” [6]. *This specifies the BMC sensor-ID namespace, not the active core count*—the distinction is the entry point for this study.

## III. METHOD

### A. Recovering the disable map

For each month we enumerate the `metric=pX_coreY_temp` partitions and count samples per (node, socket, core). A core index is *present* for a socket if it emits any sample in the window. We define the disable map of socket  $s$  as  $M[s, k] \in \{0, 1\}^{12}$  over slices  $k = \lfloor \text{core}/2 \rfloor$ .

The key validity argument is that **the BMC does not renumber active cores compactly**: were that so, every socket would report indices 0–15. Instead we observe arbitrary 16-element subsets of  $\{0, \dots, 23\}$ , so the index is a fixed physical-position identifier.

### B. Null model for clustering

Every socket has exactly four harvested slices, so the constraint induces negative dependence and a naive independence null is invalid. We use the curveball algorithm [7], which preserves *both* row sums (4 per socket) and column sums (per-slice marginals) while randomising association structure. Row- and column-sum preservation is asserted on each of 200 draws.

### C. Thermal probing of die topology

Cores on a socket share workload and inlet air, producing a dominant common mode. We remove it by subtracting the per-timestamp cross-core mean, double-centre, and compute residual correlations. Physically abutted cores should retain positive residual coupling.

*Caveat:* Common-mode removal forces each row of the residual matrix to sum to approximately zero, so negative values at large index gaps are partly mechanical. We therefore rely on contrasts *at identical index distance*, which control for this exactly.

## IV. RESULTS

### A. Every socket is 16 cores; granularity is the slice

TABLE I  
DISABLE-MAP STRUCTURE, 2020-04.

|                                               |                     |
|-----------------------------------------------|---------------------|
| Nodes analysed                                | 980                 |
| Sockets                                       | 1,960               |
| Sockets reporting exactly 16 cores            | 1,960 (100%)        |
| Sockets reporting 24 cores                    | 0                   |
| Absent cores form complete $(2k, 2k+1)$ pairs | <b>1,960 (100%)</b> |

The perfect pair structure is the central mechanical finding: it matches the POWER9 slice definition exactly and rules out per-core fusing.

### B. The map is per-die, not per-SKU

442 distinct patterns appear among the  $\binom{12}{4} = 495$  possible (Fig. 7). The most common (slices 0–3) covers only 5.7% of sockets. Within a node the two sockets match in just 1.2% of cases, with mean overlap 1.648 slices against a shuffled null of  $1.474 \pm 0.026$  ( $z = +6.6$ )—nearly independent, with a small excess attributable to same-lot pairing.

### C. Marginals are strongly non-uniform

Slice 0 is harvested on 61.7% of sockets against a uniform expectation of 33.3% ( $\chi^2 = 836.5$ ,  $df = 11$ ,  $p \ll 0.001$ ); Fig. 2 splits this by lot.

### D. Harvested slices cluster in index space

Under the curveball null the mean index gap between harvested slices is 4.243 versus  $4.585 \pm 0.016$  ( $z = -20.8$ ), and every significant co-harvest excess has  $|\Delta k| \leq 2$  (Fig. 3): slices 0&1 ( $z = +19.5$ ), 6&7 (+9.5), 8&9 (+8.7), 2&3 (+8.4), 4&5 (+7.9), 1&2 (+7.4), 10&11 (+7.1). All negative terms are distant pairs.

### E. Slice siblings are physically abutted

Fig. 4 shows residual correlation against core index distance. At *identical* distance of 1, within-slice siblings reach  $r = +0.51$  versus  $+0.29$  for cross-slice neighbours; per-node the contrast between near ( $\leq 2$ ) and far ( $\geq 12$ ) pairs is  $+0.45 \pm 0.20$  with 40/40 nodes positive (paired  $t = +13.9$ ). The two cores of a slice are roughly twice as coupled as equally-numbered cores across a slice boundary, confirming both physical abutment and the slice as a physical unit.

### F. The slice index is not established as the physical ordering

Free classical MDS on the pooled  $24 \times 24$  matrix is dominated by noise (leading axis 20.5% of variance, correlation with core index only  $+0.34$ ), so we instead score explicit candidate floorplans by correlating predicted physical distance against measured thermal coupling (Fig. 5). The identity  $1 \times 12$  linear arrangement scores  $-0.271$ , better than a random ordering ( $z = -2.22$ ,  $p = 0.025$ ) but *worse* than a  $3 \times 4$  grid ( $-0.565$ ); a hill-climbing search finds an ordering scoring  $-0.715$  whose rank correlation with the identity ordering is only  $+0.34$ .

**We therefore do not claim the sensor index is monotonic in physical die position.** The index-space clustering of Fig. 3 is a robust empirical fact, but its interpretation as *spatial* defect clustering is not established by our data and must be stated as a hypothesis. **[TODO: obtain an annotated POWER9 die shot or OpenBMC sensor-ID mapping to settle this.]**

### G. A procurement-lot boundary at node 440

Node IDs map to position as  $\text{rack} = \lfloor \text{id}/20 \rfloor$  [6]. A changepoint scan maximises Welch  $t$  at rack 22 (node 440),  $t = 27.3$  (Fig. 6).

TABLE II  
TWO POPULATIONS EITHER SIDE OF THE LOT BOUNDARY.

| Population          | Sockets | $P(\text{slice } 0)$ | Distinct patterns |
|---------------------|---------|----------------------|-------------------|
| Lot A (racks 0–21)  | 880     | 89.1%                | 224               |
| Lot B (racks 22–48) | 1,080   | 39.4%                | 414               |

Lot A marginals fall steeply with slice index; Lot B is nearly flat. Permutation tests confirm between-rack heterogeneity for slices 0, 1, 2, 5 ( $p < 0.01$ ), and same-rack sockets share more harvested slices than chance ( $1.568$  vs  $1.454 \pm 0.008$ ,  $z = +14.1$ ).

### H. Configuration lifetime

POWER9 exposes mechanisms that *could* alter the active set after manufacture: predictive Dynamic Processor Deallocation, persistent GARD records read by hostboot at IPL and clearable with `opal-gard`, and the field core override used for licensing [8], [9]. Operating-system core offlining is invisible to the BMC and cannot affect sensor presence. Whether harvest maps are static in practice is therefore an empirical question, which we answer by aggregating every core sample to (node, socket, day) counts across the dataset. **[TODO: numbers below are over the first 6 months (2020-03..2020-08); refresh once the 31-month sweep completes.]**

a) *The count is invariant:* 88.8% of socket-days report exactly 16 active cores. The remainder is dominated by *incomplete collection*, not by configuration change, on three independent grounds. (i) Reduced counts are overwhelmingly 2–10, whereas core guard removes one slice at a time and would make 14 the most common reduced value; 14 is instead among the rarest (6 socket-days). (ii) Odd counts occur (15 and 17 cores, 11 socket-days) and *cannot* arise from slice-granular fusing. (iii) Of 11,459 fully-sampled socket-days




Fig. 1. POWER9 slice organisation (after [2]) annotated with the measured per-slice harvest rate across 1,960 Marconi100 sockets. Each slice holds two SMT4 cores sharing 512kB L2 and 10MB L3 on a common fabric port. Harvesting always removes whole slices. **The horizontal ordering is the IPMI sensor index, which we show is *not* established as the physical die ordering (Sec. IV-F).**

reporting fewer than 16 cores, **99.58% report a strict subset of that socket’s modal 16**—the missing cores were simply not collected. **We find no evidence of any core-guard event in the fleet.**

*b) Configurations are static:* Restricting to *clean* socket-days (16 cores, each at >90% of the 4,320 samples expected at 20s cadence) leaves 199,135 socket-days. **94.4% of sockets never change configuration**; 106 have two distinct sets and 4 have three.

*c) Changes never occur during operation:* Of 114 transitions, **zero** occur between consecutive reporting days: every one coincides with a gap in reporting (median 2d), against a background in which 93.4% of consecutive reporting days are one day apart. Configuration changes are boot-time events, exactly as the guard/deconfiguration mechanism predicts.

*d) Changes concentrate at commissioning:* 75% of all transitions fall in 2020-03-19–21, during machine installation. In steady state ( $\geq 2020-04-01$ ) only 23 transitions occur across 16 of 981 nodes.

*e) Some apparent changes are socket relabelling:* When node  $n$ ’s socket  $p_0$  loses exactly the slices  $p_1$  gains and vice versa, the dies are unchanged and only the  $p_0/p_1$  labels have swapped. This accounts for 35% of steady-state transitions but only 4% of commissioning ones, where processors were plausibly re-seated.

*f) Caveat:* An apparent change can arise from (a) genuine hardware reconfiguration, (b) socket relabelling, or (c) a change in the BMC sensor-ID mapping, e.g. after a firmware update. The tight clustering of 86 transitions into three days is consistent with (c) and we cannot separate these from telemetry alone. **[TODO: seek CINECA maintenance logs for March 2020.]**

## V. RELATED WORK

**[TODO: expand; the citations below are located but un-verified.]** Manufacturing variability in HPC has been studied




Fig. 2. Per-slice harvest probability by procurement lot.

mainly through its *performance* consequences—frequency and power spread under a power cap [10], turbo-boost variation across nominally identical processors [11], and accelerator-level variability in GPU-rich systems [12]. That literature treats the die as a black box characterised by its behaviour. Yield modelling, conversely, treats defect statistics directly: the negative-binomial family [13] and spatial wafer-map models [14] describe clustered killer defects but are validated on fab data unavailable to system owners. To our knowledge no prior work recovers per-die harvesting maps from deployed-system telemetry, which is the gap this paper addresses.

## VI. THREATS TO VALIDITY AND OPEN ISSUES

Every unresolved item is tracked with an identifier in `OPEN_ISSUES.md` in the artefact repository, which records what is known, what would settle it, and where to look. **The items marked *blocking* below must be resolved before this paper is submitted.**

- 1) **Sensor presence  $\neq$  core existence.** Mitigated by 100% consistency across 1,960 sockets, agreement with the




Fig. 3. Co-harvest  $z$  against the curveball null. Excess is confined to near-diagonal (index-adjacent) slice pairs.




Fig. 4. Residual thermal correlation versus core index distance. The within-versus cross-slice contrast at distance 1 controls exactly for index gap.

published 16-core SKU [4], and the perfect slice-pair structure. Further supported by Sec. IV-H: sub-16 counts are collection dropout.

- 2) **[OI-1, blocking]** The sensor index may not be the physical die position (Sec. IV-F). The spatial reading of the clustering in Sec. IV-D is a hypothesis, not a result. Needs an annotated die shot or the hostboot sensor-ID mapping.
- 3) **[OI-2, blocking]** A March 2020 BMC firmware update could have remapped sensor IDs, mimicking reconfiguration; 86 of 114 transitions fall in a three-day window. Needs CINECA maintenance logs. Does not affect the steady-state or boot-time-only conclusions.
- 4) **[OI-4, blocking]** Five references are unverified and a novelty check against prior telemetry-based binning inference is outstanding.
- 5) **[OI-3]** Defect-driven versus policy-driven harvesting cannot be separated from one machine; both are reported.
- 6) **[OI-5]** Collection dropout may not be random. Monthly coverage ranges from 18% to 99.4%; if dropout correlates with lot, Sec. IV-G needs a robustness pass.
- 7) **[OI-6, OI-7]** Two small anomalies remain unex-




Fig. 5. Candidate floorplans scored against measured thermal coupling. The linear index ordering (red) beats a random ordering only marginally and is outperformed by a  $3 \times 4$  grid.




Fig. 6. Harvest probability by rack and slice. The lot boundary at rack 22 is visible as an abrupt change.

plained: 48 socket-days reporting a core outside the socket’s modal set, and two nodes whose transitions span only a one-day gap.

- 8) **[OI-8]** The mechanism behind p0/p1 relabelling is unknown; any per-socket longitudinal study on this dataset risks silently mixing two dies.
- 9) **Common-mode artefact** in the thermal analysis, addressed by same-distance contrasts; the pooled estimator may be improvable [OI-11].
- 10) **[OI-10]** Single system; findings may reflect CINECA procurement rather than POWER9 generally.

## VII. CONCLUSION

**[TODO: write.]**

## REFERENCES

- [1] A. Borghesi, C. Di Santi, M. Molan, M. S. Ardebili, A. Mauri, M. Guarrasi, D. Galetti, M. Cestari, F. Barchi, L. Benini, F. Beneventi, and A. Bartolini, “M100 ExaData: a data collection campaign on the CINECA’s Marconi100 Tier-0 supercomputer,” *Scientific Data*, vol. 10, no. 1, p. 288, 2023.
- [2] B. Thompto, “POWER9: Processor for the cognitive era,” in *Hot Chips 28 Symposium (HCS)*. IEEE, 2016, slide 11 documents 24 cores organised as 12 slices, each with 10 MB L3, attached to a 7 TB/s on-chip fabric at 256 GB/s  $\times$  12.




Fig. 7. Rank–frequency of disable patterns. Lot B is markedly more diverse.

- [3] “POWER9 — microarchitectures — IBM,” WikiChip, 2024, <https://en.wikichip.org/wiki/ibm/microarchitectures/power9; slice = 2 SMT4 cores + 512 kB L2 + 10 MB L3; die 695 mm<sup>2</sup>, 14 nm.>
- [4] IBM, “IBM Power System AC922 data sheet,” IBM Corporation, Tech. Rep., 2019, model 8335-GTG: 2 sockets, 16-core POWER9 at 2.6/3.1 GHz, 4× NVIDIA V100.
- [5] I. Baccarelli, “Introduction to the CINECA Marconi100 HPC system,” CINECA HPC user training, 2020, <https://indico.euro-fusion.org/event/341/attachments/293/664/M100.pdf>.
- [6] M100 ExaData contributors, “M100 ExaData reference repository: plugin documentation and rack layout,” [documentation/plugins/ipmi.md](https://github.com/ExaData/documentation/plugins/ipmi.md), [documentation/racks\\_spatial\\_distribution.md](https://github.com/ExaData/documentation/racks_spatial_distribution.md), 2023.
- [7] G. Strona, D. Nappo, F. Boccacci, S. Fattorini, and J. San-Miguel-Ayanz, “A fast and unbiased procedure to randomize ecological binary matrices with fixed row and column totals,” *Nature Communications*, vol. 5, p. 4114, 2014.
- [8] IBM, “POWER9: Deconfiguring hardware, examples of the reason processor cores were unconfigured, and field core override function overview,” IBM Documentation, 2021, <https://www.ibm.com/docs/en/power9?topic=configuration-deconfiguring-hardware>; covers predictive/functional deconfiguration, Dynamic Processor Deallocation, and field core override (feature code 2319).
- [9] OpenPOWER, “opal-gard: manipulate the GUARD partition on OpenPower hardware,” skiboot, [external/gard/opal-gard.1](https://github.com/open-power/skiboot), 2023, <https://github.com/open-power/skiboot>; GARD records live in a PNOR flash partition and support `list`, `show` and `clear`.
- [10] Y. Inadomi *et al.*, “Analyzing and mitigating the impact of manufacturing variability in power-constrained supercomputing,” in *SC '15*, 2015, uNVERIFIED — confirm author list and pages.
- [11] B. Acun, P. Miller, and L. V. Kale, “Variation among processors under turbo boost in HPC systems,” in *ICS '16*, 2016, uNVERIFIED — confirm pages.
- [12] “Not all GPUs are created equal: characterizing variability in large-scale, accelerator-rich systems,” *arXiv preprint arXiv:2208.11035*, 2022, uNVERIFIED — confirm author list.
- [13] C. H. Stapper, “A unified negative-binomial distribution for yield analysis of defect-tolerant circuits,” *IEEE Transactions on Semiconductor Manufacturing*, 1993, uNVERIFIED — confirm volume/pages.
- [14] “Yield prediction via spatial modeling of clustered defect counts across a wafer map,” *IIE Transactions*, 2007, uNVERIFIED — confirm authors.